

Bank Marketing (Campaign) Project Details

Team member's details

Name: Batuhan YILMAZ
Email: batuhanyilmaz1999@hotmail.com
Country: Turkey
Company: A university in Turkey
Specialization : Data Science

Group Name

Datarpher

Problem Description

ABC Bank wants to sell its term deposit product to customers and before launching the product they want to develop a model which helps them in understanding whether a particular customer will buy their product or not (based on customer's past interaction with bank or other Financial Institution).

Business Understanding

Bank wants to use ML model to shortlist customers whose chances of buying the product are more so that their marketing channel (tele marketing, SMS/email marketing etc) can focus only on those customers whose chances of buying the product are more.

This will save resources and their time (which is directly involved in the cost (resource billing)).

Project lifecycle

- **Business Understanding** (week 1)
- **Data understanding** (week 1)
- **Exploratory data Analysis** (week 2)
- **Data Preparation** (week 3)
- **Model Building (Logistic Regression, ensemble, Boosting etc)** (week 4)
- **Model Selection** (week 5)
- **Performance reporting** (week 6)
- **Deploy the model** (week 6)
- **Converting ML metrics into Business metric and explaining result to business** (week 7)
- **Prepare presentation for non technical persons.** (week 7)

Data Intake Report

Name: Bank Marketing (Campaign)

Report date: 19.07.2023

Internship Batch: LISUM22

Version:

Data intake by: Batuhan YILMAZ

Data intake reviewer: Batuhan YILMAZ

Data storage location: <https://archive.ics.uci.edu/dataset/222/bank+marketing>

Github repository link: <https://github.com/Batuhan-Ylmz/Bank-Marketing-Campaign-Term-Deposit-Product-Purchase-Classification>

Tabular data details: bank-full

Total number of observations	45211
Total number of files	1
Total number of features	17
Base format of the file	csv
Size of the data	5.9+ MB

Tabular data details: bank-additional-full

Total number of observations	41188
Total number of files	1
Total number of features	21
Base format of the file	csv
Size of the data	6.6+ MB

Data files are same as each other. However, bank-additional-full.csv file includes more specific details regarding the customers. (E.g; contact day with customer, consumer price index, consumer confidence index ...).

Data and Business Understanding:

The data is related with direct marketing campaigns of a Portuguese banking institution. The marketing campaigns were based on phone calls. Often, more than one contact to the same client was required, in order to access if the product (bank term deposit) would be ('yes') or not ('no') subscribed. The classification goal is to predict if the client will subscribe (yes/no) a term deposit (variable y).

Since there are some unnecessary kind of data that is likely to lower the quality of prediction model in the dataset with 21 features, the one dataset with 17 features (16 feature + 1 output feature) is selected as the data for model creation.

Each of 16 feature corresponds to a specific feature of that customer (feature names can be modified for easier reading later on):

- 1- age (“**age**”)
- 2- job (“**job**”)
- 3- marital status (“**marital**”)
- 4- education level (“**education**”)
- 5- if customer has credit in default (“**default**”)
- 6- yearly balance in euro (“**balance**”)
- 7- if customer has housing loan (“**housing**”)
- 8- if customer has personal loan (“**loan**”)
- 9- how did the bank contacted the customer (“**contact**”)
- 10- last contact day of the month (“**day**”)
- 11- Last contact month of the year (“**month**”)
- 12- Last contact duration (“**duration**”) → This feature highly affects the output. So, 2 models will be created with one including this feature and the other one not.
- 13- How many times customer was contacted for this campaign (“**campaign**”)
- 14- How many days have passed since the customer was contacted for the previous campaign (“**pdays**”)
- 15- Total number of contacts performed before this campaign (“**previous**”)
- 16- Did the customer subscribed for the previous term deposit product (“**poutcome**”)
- 17- Has the client subscribed a term deposit (“**y**”)

Balance of the dataset:

- The output of the data is **unbalanced**. While the ratio of people who have subscribed to the deposit product **is about 11% (“yes”)** , people who have not subscribed to the deposit product is about **%89 (“no”)**.
- Machine learning algorithms learn with the assumption that distribution of the provided labelled (output) data is symmetrical.
- In case of an unbalanced dataset, in order machine not to learn one class poorly compared to other one, **some handling methods will be applied during the EDA and Preprocessing stages**.

Missing Values:

Dataset was splitted into 2 datasets as numerical and categorical features.

- Both dataset were investigated and it was seen that there were **no missing values in the entire dataset**.
- However, some categorical columns have “**unknown**” values. Since sometimes customers with unknown features are likely to subscribe as well, they were not removed from the dataset.

Outliers:

- Numerical dataset were summarized using the 5 number summary (**min, 25%, median, 75%, max**) in order to see the distribution of each feature.
- Impact of current feature on the output was visualized to see how much the outliers influences it.
- The correlation between the numerical features and the output was investigated with a **heat-map**.
- **Observations have shown that:**
 - o *“pdays” and “previous” columns have quite less impact upon the output.*
 - o *Their variance are pretty less and consist of single values.*
 - o *People with the age range 20-60 are more likely to subscribe the deposit product.*
 - o *Duration feature has a strong impact on the output. For a healthy prediction model, 2 model will be created where first the one includes it and the second one does not.*
 - o *There is a positive skewness mostly in remaining columns.*

Approach to the outliers

- First, “**pdays**” and “**previous**” columns are dropped.
- Then, 4 different methods were picked in order to handle outlier data in the numeric_data dataset.
- As first method **Interquartile Range method (IQR)** was used.
 - o Since there is always a certain range for the remaining numerical features, coming up with an approach that eliminates the outliers that fall beyond this acceptable range can be useful.
- As second method “**winsorization**” method which is an approach based on the percentile for specific ranges was chosen.
 - o A specific percentile is selected (eg. %90, then data points smaller than %5 and greater than %95 of whole data are considered as outlier and replaced with the nearest data points).
- As third method “**MinMax Scaling**” which is an approach similar to Z-score was chosen.
 - o Rather than using the mean for outlier detection in calculation, it uses median to set the limits. Then, scales the data to a specific range. (Mostly [0,1]).
- As fourth method “**Z-score (Standard Scaling)** ” which is an approach based on re-scaling the data points that fall beyond the calculated z-score.
 - o Z-score is calculated by subtracting the current data point from the mean value and dividing it by std value for each data points

- 4 different numerical datasets were generated with the name of related outlier handling methods:
 - **Numeric_data_without_outliers_IQR,**
 - **Numeric_data_without_outliers_winsorization,**
 - **Numeric_data_without_outliers_MinMaxScaling,**
 - **Numeric_data_without_outliers_StandardScaling,**
- All these numeric datasets will be separately combined with the categorical data dataset that was created in the beginning while one version of the combined dataset having the “duration” column and the other version does not.

Hypotheses Creation

Along with the observations and exploring done, following null hypothesis (H_0) were created and all null hypotheses were **invalidated** with the calculation of ‘**z_test, p_value and chi2_contingency**’ and **alternative (H1) hypotheses were accepted:**

- **1a) Null Hypothesis (H0):** There is no significant difference in the likelihood of subscribing to the term deposit between customers aged 35 or younger and customers older than 35.
- **1b) Alternative Hypothesis (H1):** Customers aged 35 or younger are more likely to subscribe to the term deposit compared to customers older than 35.
- **2a) Null Hypothesis (H0):** There is no significant difference in the likelihood of subscribing to the term deposit between customers with a balance between -2000 and 6000 (inclusive) and customers with balances outside this range.
- **2b) Alternative Hypothesis (H1):** Customers with a balance between -2000 and 6000 (inclusive) are more likely to subscribe to the term deposit compared to customers with balances outside this range.
- **3a) Null Hypothesis (H0):** There is no significant difference in the likelihood of subscribing to the term deposit between customers with different marital statuses.
- **3b) Alternative Hypothesis (H1):** Marital status has a significant impact on the likelihood of subscribing to the term deposit.
- **4a) Null Hypothesis (H0):** There is no significant difference in the likelihood of subscribing to the term deposit between customers with different conversation duration ranges.
- **4b) Alternative Hypothesis (H1):** The duration of the conversation between customers and bank officials has a significant impact on the likelihood of subscribing to the term deposit, with the highest subscription rates observed within the duration range of 100 to 800 seconds.
- **5a) Null Hypothesis (H0):** There is no significant difference in the likelihood of subscribing to the term deposit between customers with different job categories.
- **5b) Alternative Hypothesis (H1):** The occupation of the customer has a significant impact on the likelihood of subscribing to the term deposit, with customers in the job categories 'management', 'technician', and 'blue-collar' showing higher subscription rates compared to other job categories.

Cleaning and Transformation of Categorical Data

- Even though there are no missing values in categorical columns, high amount of “unknown” category is exist in one of the features of categorical_data. However, this “unknown” type category makes up more than the %80 of whole data whose modification might cause huge misleads for the learning process of model. So, it is treated as a separate and valid category like the others.

Approach for Encoding the Categorical Columns

- For handling the encodings of categorical data, the most common encoding approaches will be used:
 - o **One-hot encoding**
 - o **Label Encoding**
- Since one-hot encoding creates new columns as many as the number of categories in that feature, using one-hot encoding for more than a threshold value that of the specific dataset– e.g, depending on the feature numbers in dataset this can be around 10- it may lead to a high-dimensional feature space which
 - o **Can be computationally more expensive,**
 - o **Lead to the situation known as “curse of dimensionality”.**
 - **Such as:**
 - | Education_primary | education_high_school | education_college |
|-------------------|-----------------------|-------------------|
| 0 | 1 | 0 |
| 1 | 0 | 0 |
| | | |
- In case there are too many different categories for the features, Label Encoding is used to specify an integer for the corresponding category.
 - **Such as:**
 - | Education | month |
|-----------|-------|
| 0 | 3 |
| 7 | 8 |
| 2 | 6 |
| | |
- In the categorical dataset, only the “**month**” and “**job**” columns contains more than 10 categories.
- Hence, only these 2 categorical features (“**month**”, “**job**”) were encoded using “**LabelEncoder**” in order to prevent high-dimensionality.
- Rest of the categorical features encoded using **OneHotEncoder**.

Currently there are **4 different numeric_features** without outliers processed with **4 different outlier handling algorithms** and categorical_features where features with more than 10 are encoded with LabelEncoder and less than 10 are encoded with OneHotEncoder.

Preprocessed Datasets Creation

As mentioned before, 4 different outlier handling methods were used:

- **Inter Quartile Range (IQR)**
- **Winsorization**
- **Min-Max Scaling**
- **Z- Score (Standard Scaling)**

Only **IQR** method **removes** the outlier data in dataset rather than modifying it just as the other 3 methods. Hence, whole dataset were first cleaned with the IQR outlier handling method first, only then splitted into numerical and categorical features for further preprocessing.

Then, same process is applied for all the other methods and datasets were named as:

- preprocessed_IQR.csv
- preprocessed_winsorization.csv
- preprocessed_MinMax.csv
- preprocessed_StdScal.csv

All the datasets (except for the “IQR”) have the following shape:

Datasets Name:

- preprocessed_winsorization.csv,
- preprocessed_MinMax.csv,
- preprocessed_StdScal.csv

Total number of observations	45211
Total number of files	3
Total number of features	28
Base format of the file	csv
Size of the data	3.3 MB

IQR-way-handled dataset has the following shape:

Dataset Name: preprocessed_IQR.csv

Total number of observations	34719
Total number of files	1
Total number of features	28
Base format of the file	csv
Size of the data	2.8 MB

However, as the meta-data suggest, the feature “**duration**” has a strong correlation with the target (output) feature “**y**”. **Therefore, 1 extra form of each dataset without the duration feature was created as well.**

New datasets without the “duration” feature (except for the “IQR”) have the following shape:

Datasets Name:

- preprocessed_winszORIZATION_without_duration.csv,
- preprocessed_MinMax_without_duration.csv,
- preprocessed_StdScal_without_duration.csv

Total number of observations	45211
Total number of files	3
Total number of features	27
Base format of the file	csv
Size of the data	2.9 MB

IQR-way-handled dataset without the “duration” feature has the following shape:

Dataset Name:

- preprocessed_IQR_without_duration.csv

Total number of observations	34719
Total number of files	1
Total number of features	27
Base format of the file	csv
Size of the data	2.5 MB

Therefore, total 8 different datasets with following names are on point in final form:

- 1- preprocessed_IQR.csv
- 2- preprocessed_IQR_without_duration.csv
- 3- preprocessed_MinMax.csv
- 4- preprocessed_MinMax_without_duration.csv
- 5- preprocessed_StdScal.csv
- 6- preprocessed_StdScal_without_duration.csv
- 7- preprocessed_winszORIZATION.csv
- 8- preprocessed_winszORIZATION_without_duration.csv

Unbalanced Class Handling

As discussed before :

- virtually **89 %** of total output (target) data is belong to **class 0 (“no”)**
- **11%** of total output (target) data is belong to **class 1 (“yes”)**.

This can cause models to learn with a bias towards the majority class. Hence, a poor performance for the minority class.

In order to solve this issue, a few different approaches can be considered and tried one by one:

1- Sampling

- a. **Under Sampling:** This technique randomly reduces the number of features in the dataset that are belong to majority class so that model will be prevented from being biased towards majority class.
- b. **Over Sampling:** This technique increases the number of instances in minority class by duplicating or generating new instances so that there will be more data to train with for the models regarding the minority class.

2- Adjusting Class Weights

- o Many machine learning algorithms allow class weights arrangement during the training. Weights of the minority and majority classes can be indicated manually (e.g: , in our case 8 for minority and 2 for majority class so that percentages will match) so that importance of the classes will be emphasized and model will be trained taking the weights into consideration.

3- Different Algorithms

- a. Different machine learning algorithms might perform better when it comes to unbalanced data as the way they are built functionally is relatively less-sensitive to the unbalanced data.
- b. Models such as:
 - i. **Gradient Boosting Algorithms (e.g: XGBoost, LightGBM..etc)**
 - ii. **SVR, SVM,**
 - iii. **Neural Networks (depends on how the architecture is defined),**
 - iv. **Naïve Bayes**
 - v. **Decision Trees**
 - vi. **Ensemble Methods (e.g: Random Forest)**

4- Different Evaluation Metrics

- a. Since the traditional metrics might focus on specific points when evaluating that whether the model is overfitted or underfitted for one class, they might be inadequate to full assess if any poor training is occurred.
- b. Using all kind of evaluation metrics such as
 - i. **Precision,**
 - ii. **Recall,**
 - iii. **F1-score,**
 - iv. **Area under the ROC curve (AUC-ROC)**

can give us a better idea about the model's performance on such unbalanced data.

5- Cross Validation

- a. Cross-validation helps assess how well the model will generalize to new, unseen data. In imbalanced cases, it's important to ensure that each fold maintains the same class distribution as the original dataset to avoid introducing bias.

(EDA Presentation is included to the project before model creation and fitting).